

How this was actually made

You're about to see how one question — **"should I put money into this company?"** — became a fact-checked investment analysis. Then you'll watch it get made again, live, from scratch — to see whether the process actually holds up, or just got lucky the first time.

First, the honest part most demos skip

There's a lot of money changing hands right now for "AI work" that's really a confident paragraph and a shrug. This isn't that. The polished documents you're about to see are the **easy** part — they come last and take the least time. The work that *matters* is the invisible kind: **checking every fact against two independent sources, catching its own mistakes, and flatly refusing to make anything up.** That rigor is the product; the pretty report is just its surface. Today you get to see underneath.

The three documents (a 2-minute look)

🕒 THE DOSSIER

The interactive brief

What Jimmy opens. Every section has a notes box — *his* answers are what sharpen the final version. Built to be reacted to, not just read.

🕒 THE FROZEN FIRST CUT

Where it started

An earlier version, locked and saved untouched — so we can prove how much the work deepened, and compare honestly.

🕒 THE FULL ANALYSIS

The heavyweight

The deep version built *without* any of Jimmy's input yet — every competitor profiled, every number verified twice. The "if we had to decide today" read.

The black box

Every single action the system took was stamped, in order, to a **flight recorder** — a plain log, ~120 entries, no spin. Which tool it called and whether it worked. Where it had to make a judgment call. Every time it caught its own error. A real one:

```
GATE disambiguation = PASS w/ flag: the company isn't the "AI hiring tool" it was pitched as → flag, don't build on it
REVIEW sanity-check flagged a possible hallucination in my own work → re-verified before trusting it
REVIEW independent evaluator caught a flaw I'd missed → fixed it
JUDGMENT "$30M, not $51M" — a funding figure mislabeled; corrected (fatal in a real analysis)
```

It can't spin you — it wrote down everything it did, including the ugly parts. That log is the thing nobody selling you AI can show you: proof of *how*, not just *what*.

◆ NOW THE EXPERIMENT — AND WE DIDN'T RIG IT

A month ago, before any of this was built, we sealed a written **prediction** of how a fresh AI would tackle this job — including the nine places we bet it would *cut corners*. Then we ran it carefully and logged the ideal path. We never tested the prediction honestly... until right now.

1

Sean opens a **brand-new session** — no memory, clean slate.

2

He pastes in the **original prompt** — the same starting question, nothing added.

3

We watch it run **live**, and tick off, in real time: did it reproduce the same rigor — or diverge?

Rigged demos are easy; this isn't one. Either it proves the process is real and repeatable, or it shows us exactly where it isn't — and we leave smarter either way. The boring, check-it-twice stuff is what holds up when there's real money on the table. Turns out that's the hard part to fake.

Ready?